

# CUSTOMER CHURN PREDICTION

## Problem Definition and Design Thinking

In this part you will need to understand the problem statement and create a document on what have you understood and how will you proceed ahead with solving the problem.

Please think on a design and present in form of a document.

**Project Definition:** The project involves using IBM Cognos to predict customer churn and identify factors influencing customer retention. The goal is to help businesses reduce customer attrition by understanding the patterns and reasons behind customers leaving. This project includes defining analysis objectives, collecting customer data, designing relevant visualizations in IBM Cognos, and building a predictive model.

## Design Thinking:

1. **Analysis Objectives:** Define the specific objectives of predicting customer churn, such as identifying potential churners and understanding the key factors contributing to churn.
2. **Data Collection:** Determine the sources and methods for collecting customer data, including customer demographics, usage behavior, and historical interactions.
3. **Visualization Strategy:** Plan how to visualize the insights using IBM Cognos, showcasing factors affecting churn and retention rates.
4. **Predictive Modeling:** Decide on the machine learning algorithms and features to use for predicting customer churn.

**Dataset Link:** <https://www.kaggle.com/datasets/blastchar/telco-customer-churn>

Creating a design to **analyse customer churn** and implementing it involves several steps, from defining the problem to implementing a solution. Here's a detailed guide on how to go about it:

**Step 1: Problem Definition and Understanding**

1. Define the Objective: Clearly state the goal of your customer churn analysis, such as predicting which customers are likely to churn.
2. Understand Churn: Gain a deep understanding of what customer churn means for your business, what factors contribute to it, and why it's important to address.
3. Gather Stakeholder Requirements: Collaborate with stakeholders (e.g., product managers, marketing teams) to understand their requirements and expectations regarding the churn analysis.

**Step 2: Data Collection and Preparation**

1. Identify Data Sources: Determine the sources of data relevant to customer churn, such as customer profiles, usage patterns, transaction history, etc.
2. Data Cleaning and Integration: Clean the data to remove errors, duplicates, and irrelevant entries. Integrate data from various sources into a cohesive dataset.
3. Feature Engineering: Identify relevant features (attributes) that might impact customer churn (e.g., customer demographics, usage patterns, satisfaction scores) and engineer new features if needed.

**Step 3: Exploratory Data Analysis (EDA)**

1. Descriptive Statistics: Analyze basic statistics of the data to understand its distribution and characteristics.
2. Visualizations: Create various plots and visualizations (e.g., histograms, scatter plots, correlation matrices) to identify patterns and relationships in the data.
3. Identify Key Factors: Use EDA to identify the factors most likely to influence customer churn.

#### **Step 4: Model Selection and Development**

1. Model Selection: Choose appropriate machine learning models for the churn prediction task (e.g., logistic regression, decision trees, random forests, neural networks).
2. Train-Test Split: Divide the dataset into training and testing sets to evaluate the model's performance.
3. Model Training: Train the chosen model using the training data, optimizing hyperparameters for best performance.

#### **Step 5: Model Evaluation**

1. Performance Metrics: Evaluate the model's performance using relevant metrics (e.g., accuracy, precision, recall, F1-score) to assess its predictive capabilities.
2. Fine-Tuning: If needed, fine-tune the model or experiment with different algorithms to improve performance.

#### **Step 6: Implementation and Deployment**

1. Model Integration: Integrate the trained and validated model into the organization's existing systems or platforms.
2. Real-Time Prediction: Set up a mechanism for real-time prediction using the deployed model to predict customer churn as new data becomes available.
3. Monitoring and Maintenance: Continuously monitor the model's performance in the live environment, retraining it periodically to ensure it remains accurate and effective.

#### **Step 7: Interpretation and Actionable Insights**

1. Interpret Results: Interpret the model's predictions and understand the factors contributing to customer churn as identified by the model.
2. Recommendations: Generate actionable insights and recommendations based on the model's findings to mitigate churn (e.g., targeted marketing strategies, product improvements).

## ANALYSIS IMPLEMENTATION:

```
#import all relevant libraries

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import missingno as msno
import warnings
warnings.filterwarnings("ignore")
from sklearn.experimental import enable_iterative_imputer
from sklearn.impute import IterativeImputer
from sklearn.preprocessing import LabelEncoder
```

```
#Loading the dataset
data=pd.read_csv("C:\Users\HP\Downloads\Telco-Customer-Churn.csv")
```

### Data Collection:

- We will collect customer churn prediction data from reputable sources pertinent to our industry, such as internal databases, customer records, and transaction histories.
- The primary data source for this project will be our company's customer database, encompassing historical customer interactions, subscription details, and churn-related information.
- Daily data updates will be drawn from our internal customer records, allowing us to maintain a real-time understanding of churn.
- Subsequently, this data will be merged with additional information, such as demographic data, to enrich our analysis and develop a more comprehensive customer churn prediction model.

```
data.head()
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	Device
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service	DSL	No	Yes	No
1	5575-GNVDE	Male	0	No	No	34	Yes	No	DSL	Yes	No	Yes
2	3668-QPYBK	Male	0	No	No	2	Yes	No	DSL	Yes	Yes	No
3	7795-CFOCW	Male	0	No	No	45	No	No phone service	DSL	Yes	No	Yes
4	9237-HQITU	Female	0	No	No	2	Yes	No	Fiber optic	No	No	No

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 7043 entries, 0 to 7042
```

```
Data columns (total 21 columns):
```

#	Column	Non-Null Count	Dtype
0	customerID	7043 non-null	object
1	gender	7043 non-null	object
2	SeniorCitizen	7043 non-null	int64
3	Partner	7043 non-null	object
4	Dependents	7043 non-null	object
5	tenure	7043 non-null	int64
6	PhoneService	7043 non-null	object
7	MultipleLines	7043 non-null	object
8	InternetService	7043 non-null	object
9	OnlineSecurity	7043 non-null	object
10	OnlineBackup	7043 non-null	object
11	DeviceProtection	7043 non-null	object
12	TechSupport	7043 non-null	object
13	StreamingTV	7043 non-null	object
14	StreamingMovies	7043 non-null	object
15	Contract	7043 non-null	object
16	PaperlessBilling	7043 non-null	object
17	PaymentMethod	7043 non-null	object
18	MonthlyCharges	7043 non-null	float64
19	TotalCharges	7043 non-null	object
20	Churn	7043 non-null	object

```
dtypes: float64(1), int64(2), object(18)
```

```
memory usage: 1.1+ MB
```

```
data.describe()
```

	SeniorCitizen	tenure	MonthlyCharges
count	7043.000000	7043.000000	7043.000000
mean	0.162147	32.371149	64.761692
std	0.368612	24.559481	30.090047
min	0.000000	0.000000	18.250000
25%	0.000000	9.000000	35.500000
50%	0.000000	29.000000	70.350000
75%	0.000000	55.000000	89.850000
max	1.000000	72.000000	118.750000

```
data[[col for col in data.columns.difference(num_cols) if col != 'seniorcitizen']]
.describe().T
```

	count	unique	top	freq
churn	7043	2	No	5174
contract	7043	3	Month-to-month	3875
dependents	7043	2	No	4933
deviceprotection	7043	3	No	3095
gender	7043	2	Male	3555
internetservice	7043	3	Fiber optic	3096
multiplelines	7043	3	No	3390
onlinebackup	7043	3	No	3088
onlinesecurity	7043	3	No	3498
paperlessbilling	7043	2	Yes	4171
partner	7043	2	No	3641
paymentmethod	7043	4	Electronic check	2365
phoneservice	7043	2	Yes	6361
streamingmovies	7043	3	No	2785
streamingtv	7043	3	No	2810

### Data Preprocessing

- Data cleaning and preprocessing constitute crucial phases in the preparation of data for analysis.
- During this stage, we will address data issues including duplicate entries, inconsistent formats, the treatment of missing values, and the conversion of categorical variables into numerical formats.

```
data.dtypes
```

```
customerID      object
gender          object
SeniorCitizen   int64
Partner         object
Dependents      object
tenure          int64
PhoneService    object
MultipleLines   object
InternetService object
OnlineSecurity  object
OnlineBackup    object
DeviceProtection object
TechSupport     object
StreamingTV     object
StreamingMovies object
Contract        object
PaperlessBilling object
PaymentMethod   object
MonthlyCharges  float64
TotalCharges    object
Churn           object
dtype: object
```

```
data.isnull().sum()
```

```
customerID      0
gender          0
SeniorCitizen   0
Partner         0
Dependents      0
tenure          0
PhoneService    0
MultipleLines   0
InternetService 0
OnlineSecurity  0
OnlineBackup    0
DeviceProtection 0
TechSupport     0
StreamingTV     0
StreamingMovies 0
Contract        0
PaperlessBilling 0
PaymentMethod   0
MonthlyCharges  0
TotalCharges    0
Churn           0
dtype: int64
```

---

Data Exploration

- Engage in exploratory data analysis (EDA) to grasp the dataset's characteristics, patterns, and relationships.
- In this phase, the objective is to delve into the data to comprehend its nuances. EDA entails computing statistical summaries, creating data visualizations, and recognizing patterns and anomalies.
- Key areas of exploration encompass customer demographics, historical usage patterns, and the impact of various features on churn predictions.
- Utilize visualizations to gain insights into the distribution of key features and the identification of influential factors affecting customer churn.

```
#data cleaning  data transformation  data reduction
#drop irrelevant variables
data=data.drop(['CustomerId'],axis=1)
#identifying and treating missing values
data.isnull().sum()
data.fillna(0)

data.head()
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	DeviceProtection	TechSupport
0	Female	0	Yes	No	1	No	No phone service	DSL	No	Yes	No	No
1	Male	0	No	No	34	Yes	No	DSL	Yes	No	Yes	No
2	Male	0	No	No	2	Yes	No	DSL	Yes	Yes	No	No
3	Male	0	No	No	45	No	No phone service	DSL	Yes	No	Yes	Yes
4	Female	0	No	No	2	Yes	No	Fiber optic	No	No	No	No



```

num_cols = ['tenure', 'monthlycharges', 'totalcharges']

label="Churn"

plt.figure(figsize = (15, 26))

for i, col in enumerate(data.columns.difference(num_cols)[1:]):plt.subplot(6, 3, i+1)

    ax = sns.countplot(data, x = col, hue = label)

    ax.bar_label(ax.containers[0])

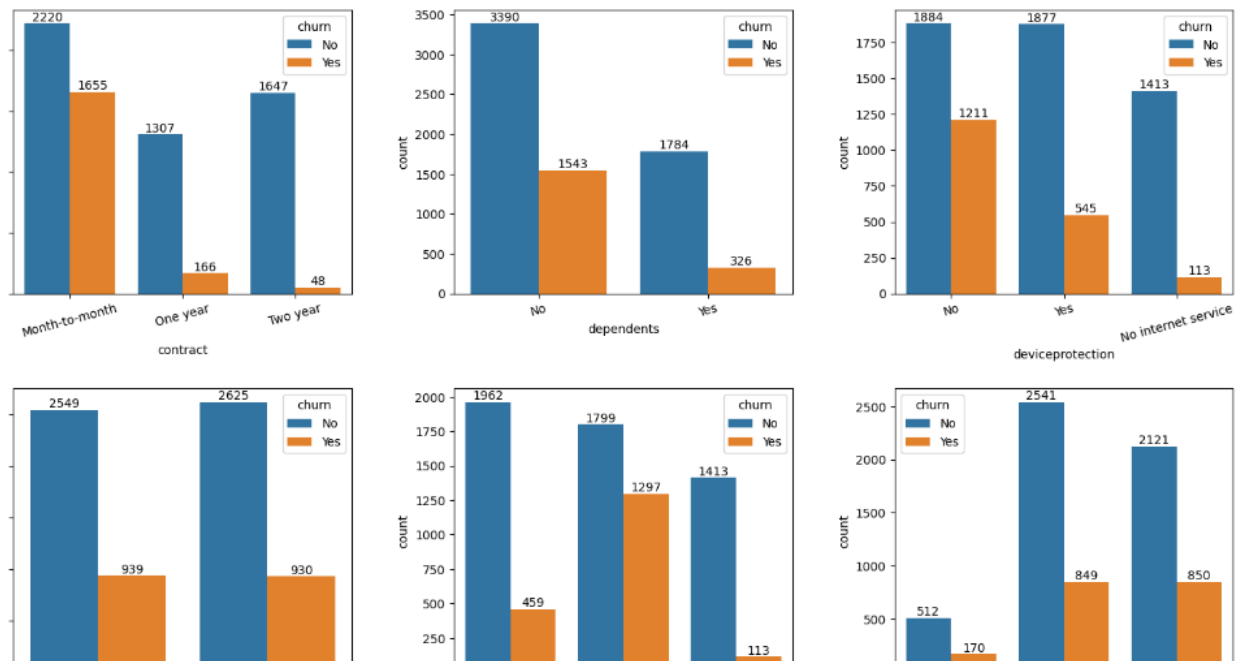
    ax.bar_label(ax.containers[1]) plt.xticks(rotation

    = 15)


plt.tight_layout()

plt.show()

```



```
plt.figure(figsize = (4,4))
ax = sns.countplot(data, x = label)
ax.bar_label(ax.containers[0])
plt.show()
```



## PREDICTIVE MODEL:

LOGISTIC REGRESSION

```
In [1]: y = df_dummies['Churn'].values
X = df_dummies.drop(columns = ['Churn'])

# Scaling all the variables to a range of 0 to 1
from sklearn.preprocessing import MinMaxScaler
features = X.columns.values
scaler = MinMaxScaler(feature_range = (0,1))
scaler.fit(X)
X = pd.DataFrame(scaler.transform(X))
X.columns = features
```

It is

important to scale the variables in logistic regression so that all of them are within arange of 0 to 1. This helped me improve the accuracy from 79.7% to 80.7%.

```
In [2]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=101)
```


```
In [ ]: from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
result = model.fit(X_train, y_train)
```

```
In [3]: from sklearn import metrics
prediction_test = model.predict(X_test)
# Print the prediction accuracy
print (metrics.accuracy_score(y_test, prediction_test))
```

0.8075829383886256

```
In [4]: # To get the weights of all the variables
weights = pd.Series(model.coef_[0],
                    index=X.columns.values)
print (weights.sort_values(ascending = False)[:10].plot(kind='bar'))
```

AxesSubplot(0.125,0.125;0.775x0.755)



```
In [5]: print(weights.sort_values(ascending = False)[-10:].plot(kind='bar'))
```

AxesSubplot(0.125,0.125;0.775x0.755)

